LIFTING THE CURSE OF DIMENSIONALITY

A DASHBOARD FOR HIGH DIMENSIONAL DATA VISUALIZATION

MACHINE LEARNING IN PRACTICE

Analysis Pipeline:

- Collect, clean, explore data
- Process data to develop features and select model
- Choose a model that uses features to answer a question

The Problem with Big Data:

- Haphazard collection. Interferes with data exploration stage: what is relevant? What questions can the dataset answer?
- Classifier performance: as features/dimensions are added, similarity become harder to measure mathematically (the distance between features becomes HUGE)

GOOD PREDICTION, HARD INTERPRETATION?

Just Use “Black Box” Methods?

- Like: Neural Networks, Random Forests
- Great for classifying/predicting
- Not good enough for science: we lose the ability to interpret the model
- Not good enough for small data: don’t work as well as hand-crafted models
- e.x.: predict patient visual stimulus using fMRI data. Good prediction with opaque model doesn’t answer questions of what part of brain is most active

PREVIOUS APPROACHES: CLASSICAL STATISTICS

Statisticians have been dealing with the data exploration problem for at least a century. Their tools:

- Scatterplot matrices
  - Indispensable for finding pairwise relationships
- Histograms
  - Expose distribution of a particular feature
- The problem: suitable only for low dimensions

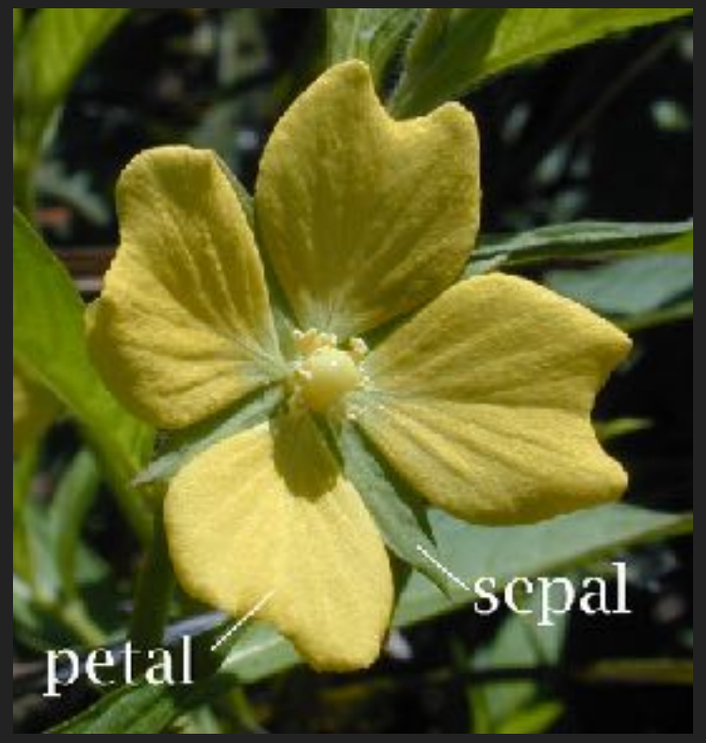

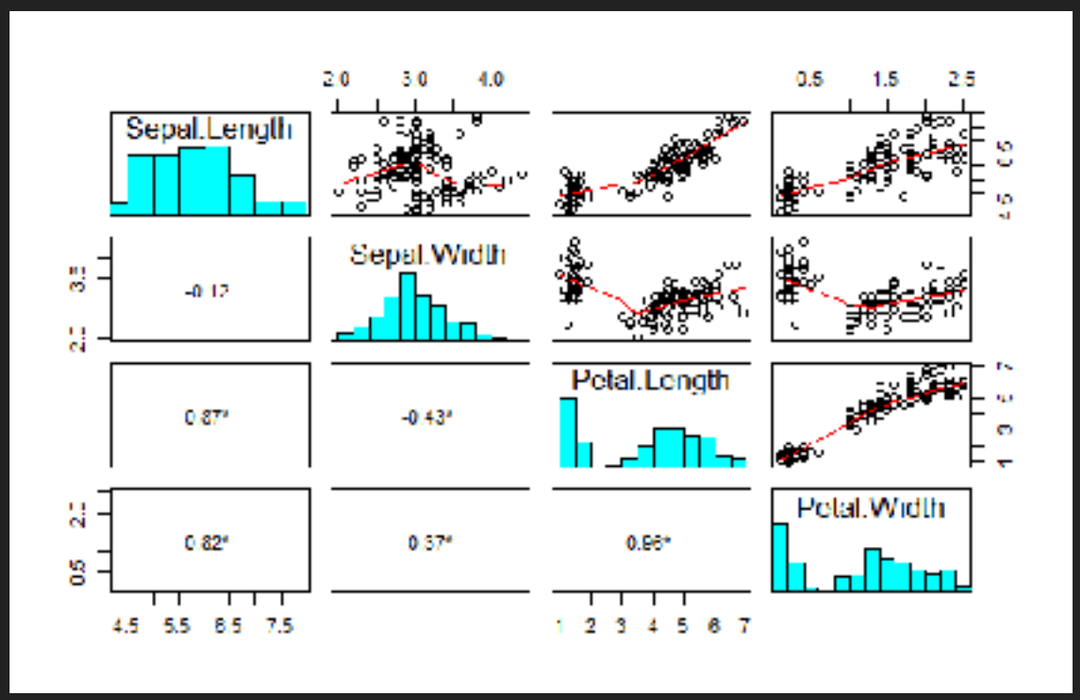

UP TO 9 DIMENSIONS, STILL USEFUL GIVEN MODIFICATIONS:

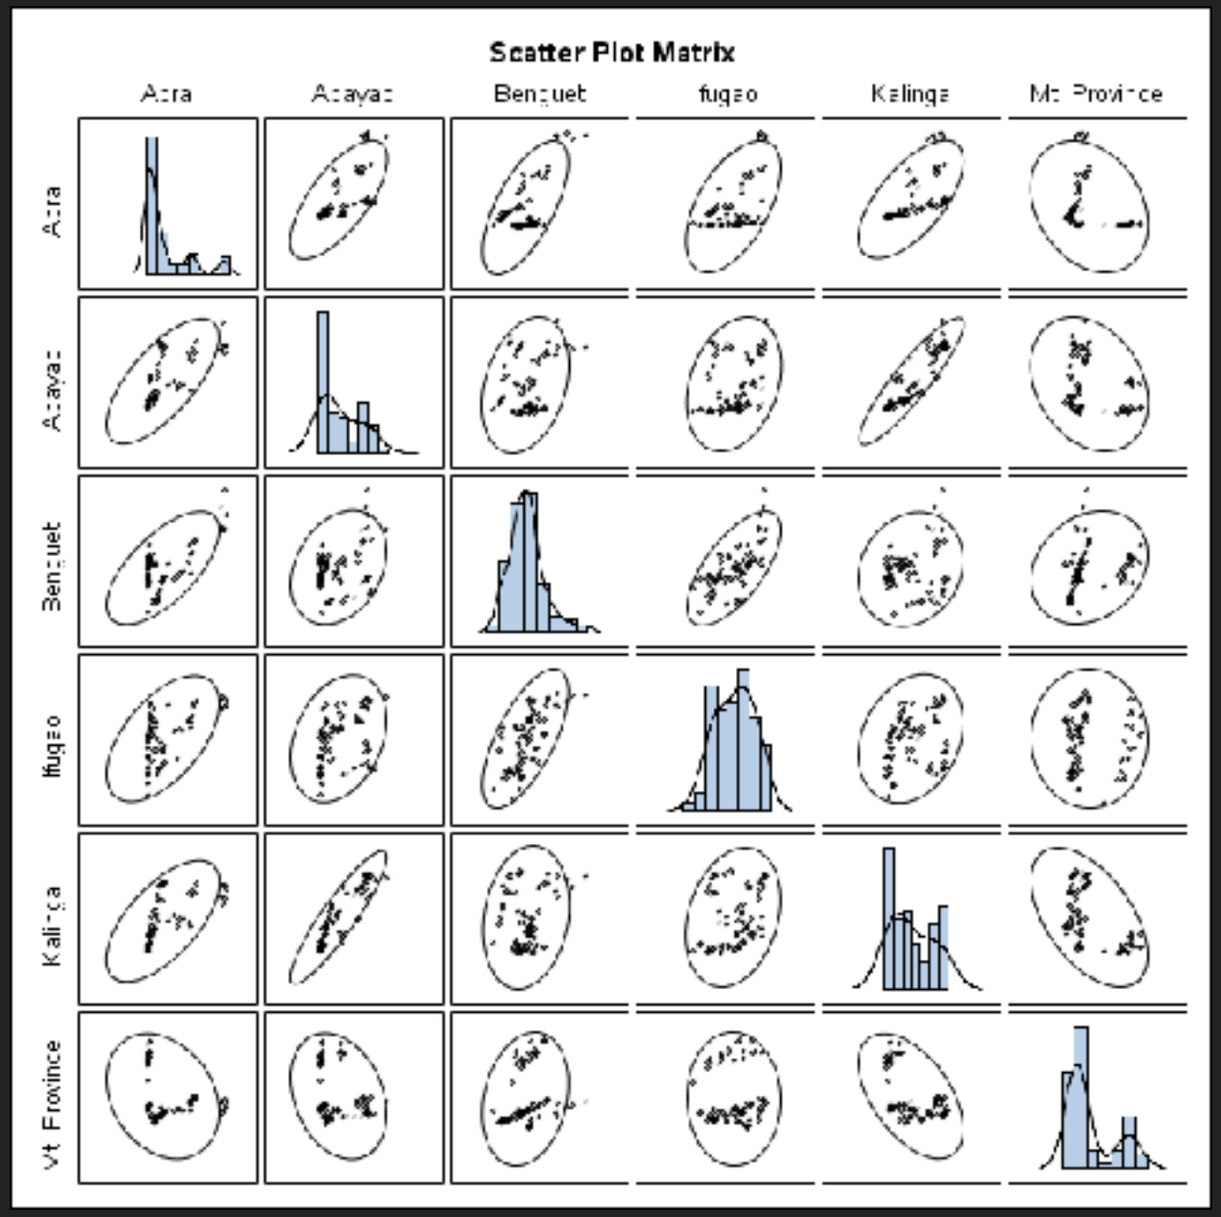

MORE THAN 9 DIMENSION, USELESS FOR EXPLORATION:

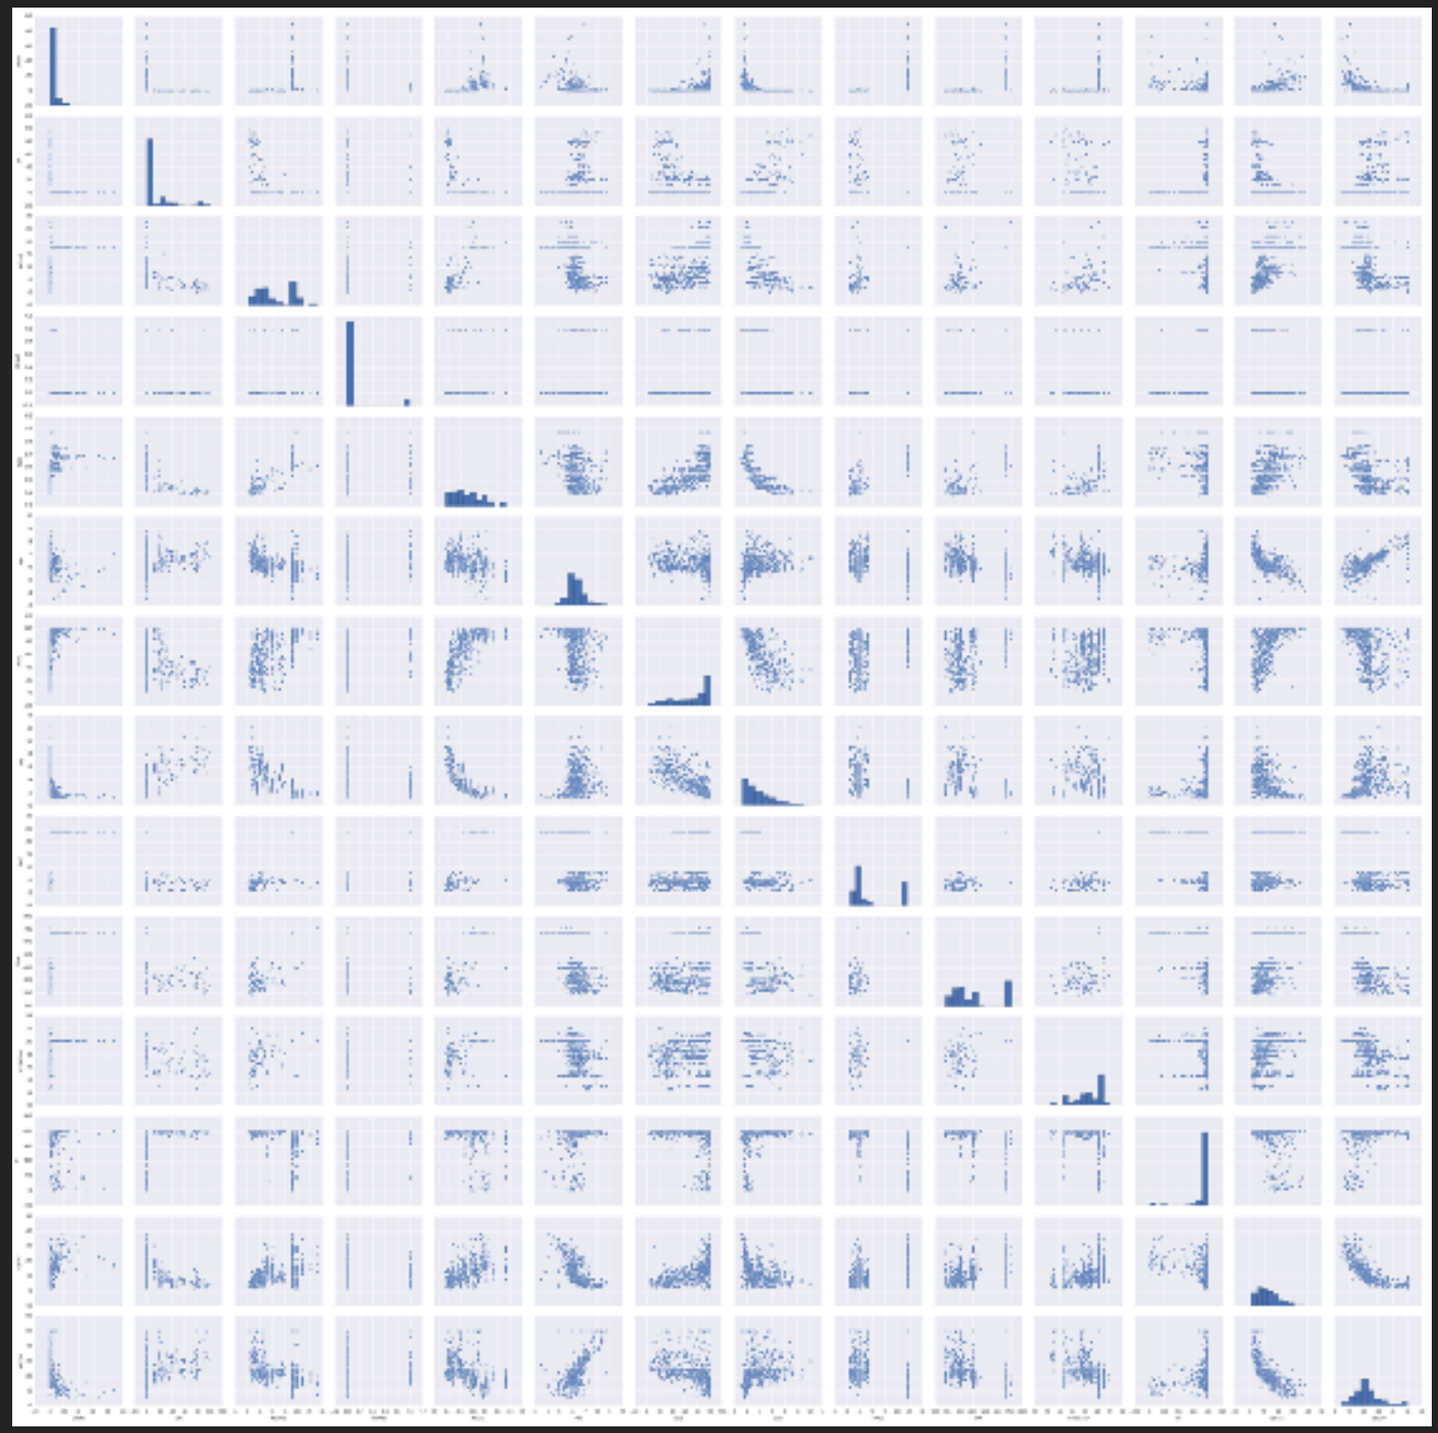

MACHINE LEARNING: BUILDING HUMAN-INTERPRETABLE MODELS

MACHINE LEARNING NEEDS NEW VISUALIZATION TOOLS

- No Free Lunch theorem
  - Model exploration is a larger concern
  - Tools like d3.js allow dynamic exploration of complex data
- But: high dimensionality of data interferes with naive dynamic visualizations
- My contribution: a visualization dashboard that combines algorithms for dimensionality reduction with visualization

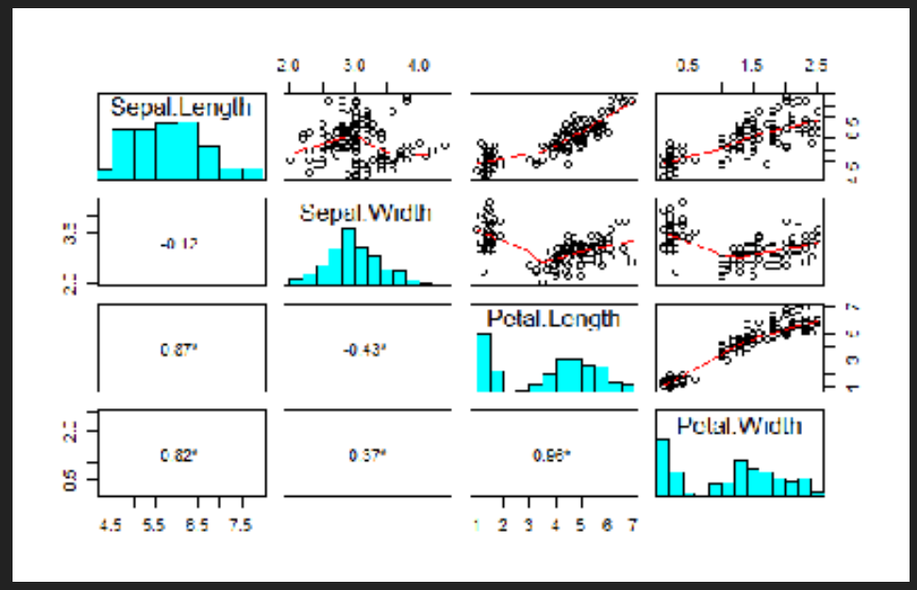

MACHINE LEARNING: BUILDING HUMAN-INTERPRETABLE MODELS

MACHINE LEARNING NEEDS NEW VISUALIZATION TOOLS

- No Free Lunch theorem
  - Model exploration is a larger concern
  - Tools like d3.js allow dynamic exploration of complex data
- But: high dimensionality of data interferes with naive dynamic visualizations
- My contribution: a visualization dashboard that combines algorithms for dimensionality reduction with visualization

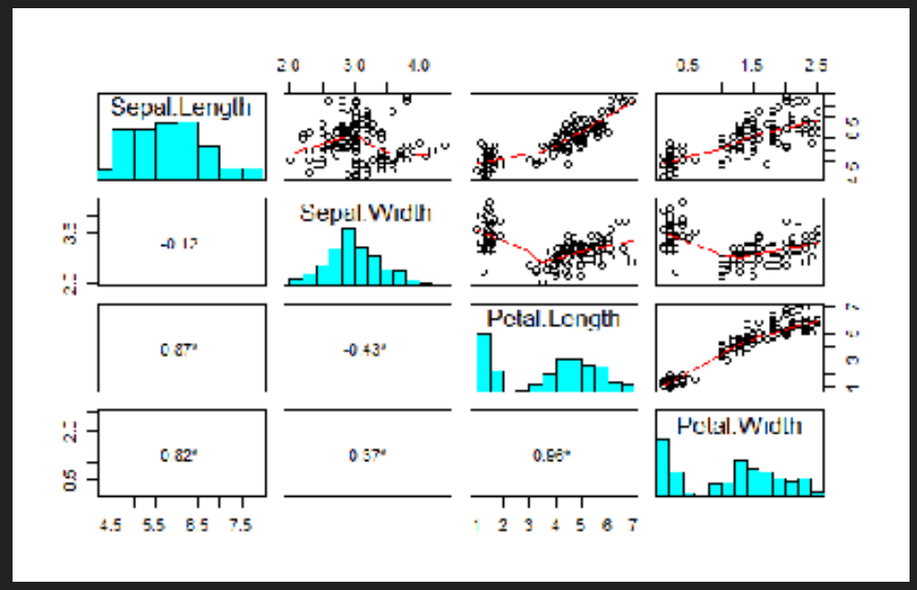

DRAG-AND-DROP DATA FILE INTO APPLET

- Each time new comma-separated file is dragged into interface, the plots are cleared
- Optional labels file for pre-labelled data
- Highlight animation when file is dragged over top
- Shows file name when loaded

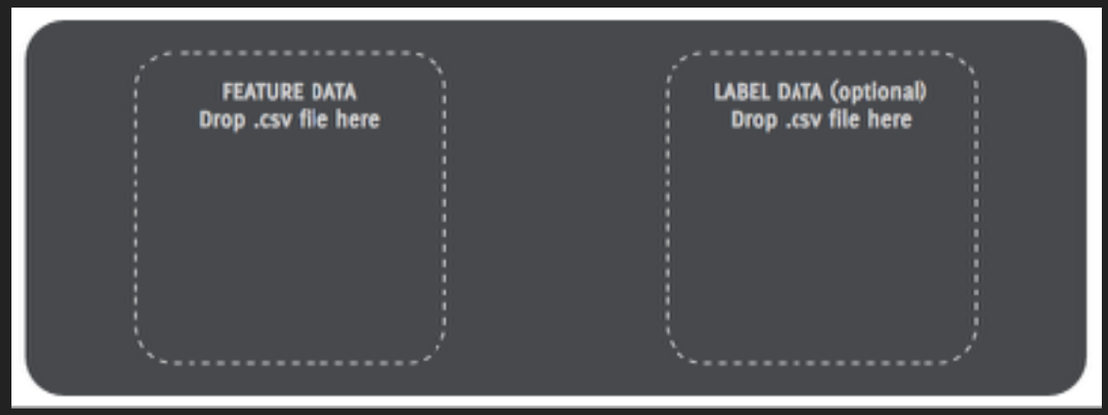

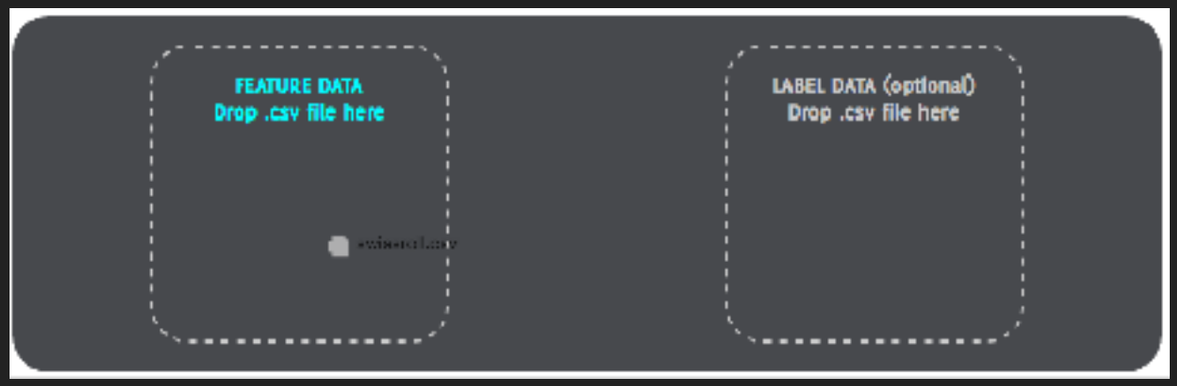

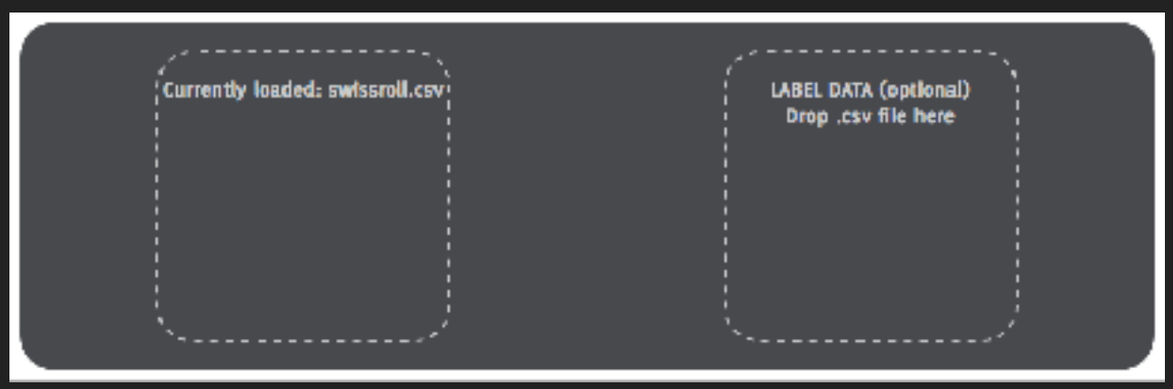

VISUALIZATION WIDGETS: REDUCE AND CLUSTER

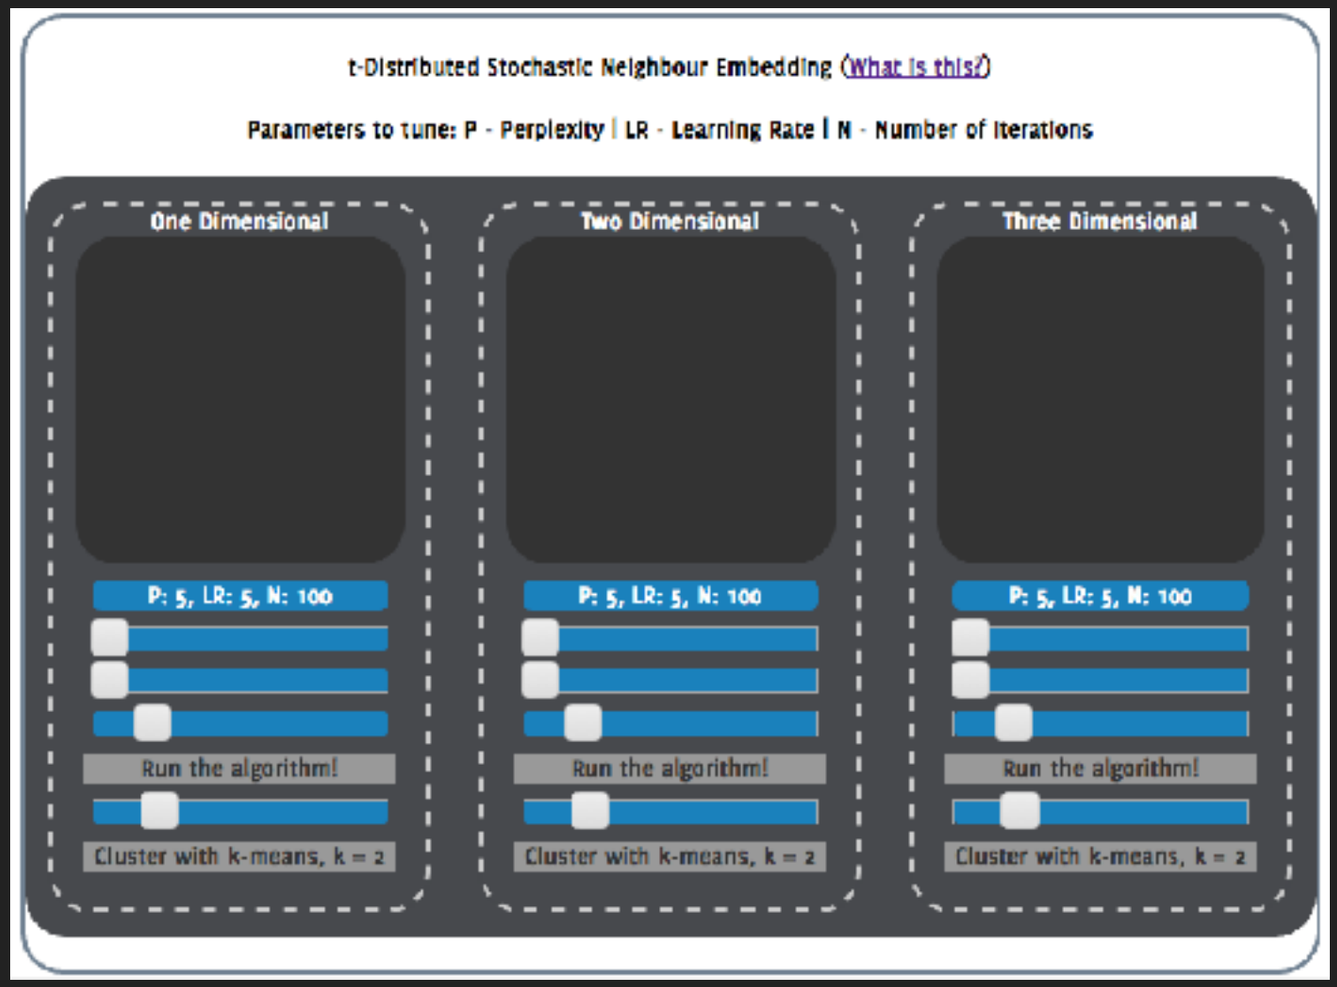

- Each dimensionality reduction algorithm has its own interface widget
- 3 divisions per algorithm for each of 3 possible visualization for reduction
- Each reduction can be tuned individually and compared against other three

VISUALIZATION WIDGET: ACTIONS

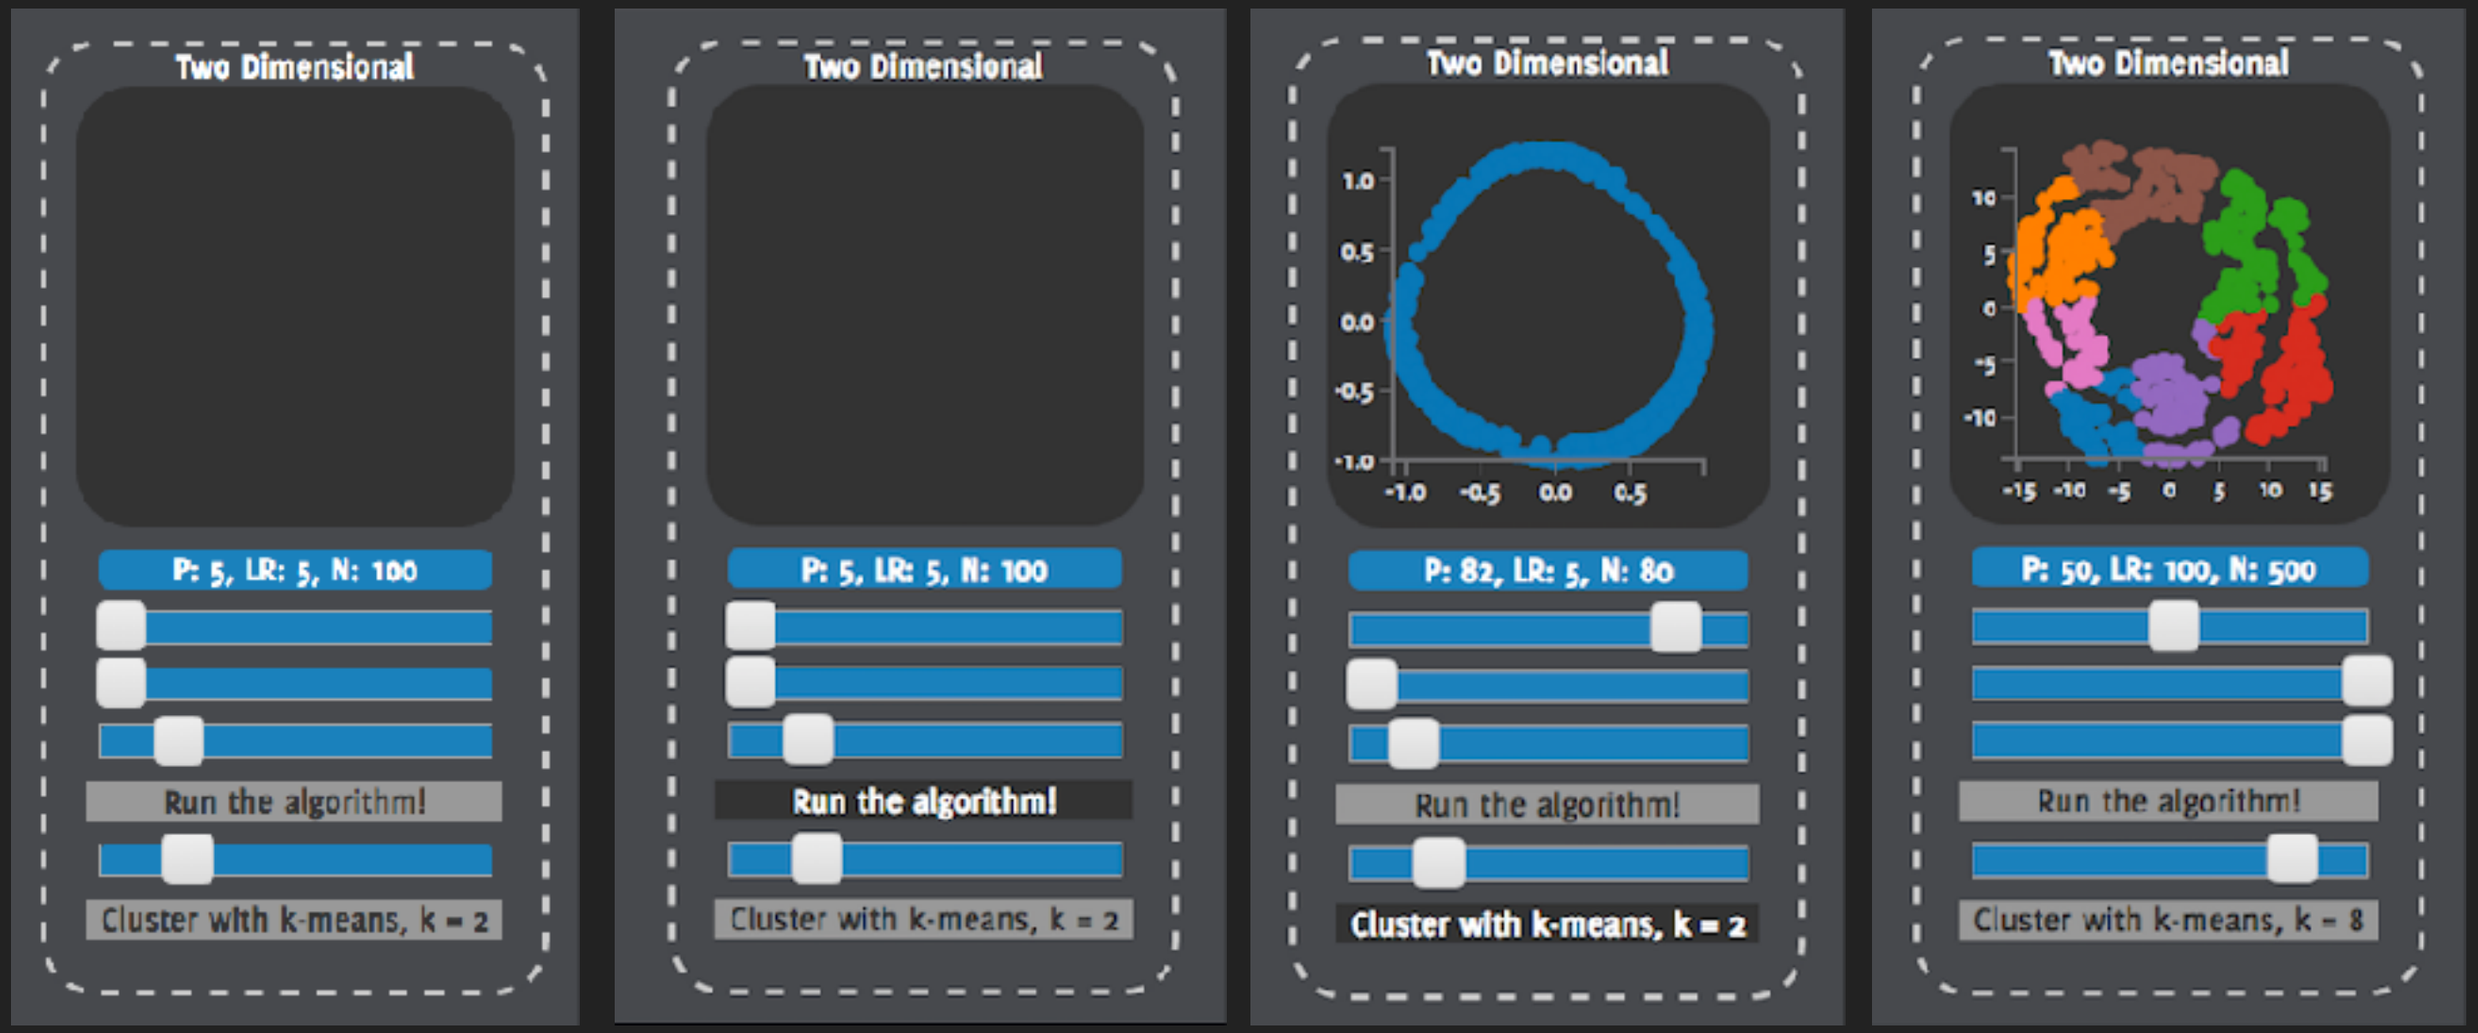

- Tunable parameters are depicted below the plotting region
- Individual sliders are used to select the values of the parameters
- Buttons invert color to show the possibility of actions
- k-means clustering is available at the bottom to find relationships using similarity by distance

VISUALIZATION WIDGET: 3D DATA EXPLORATION

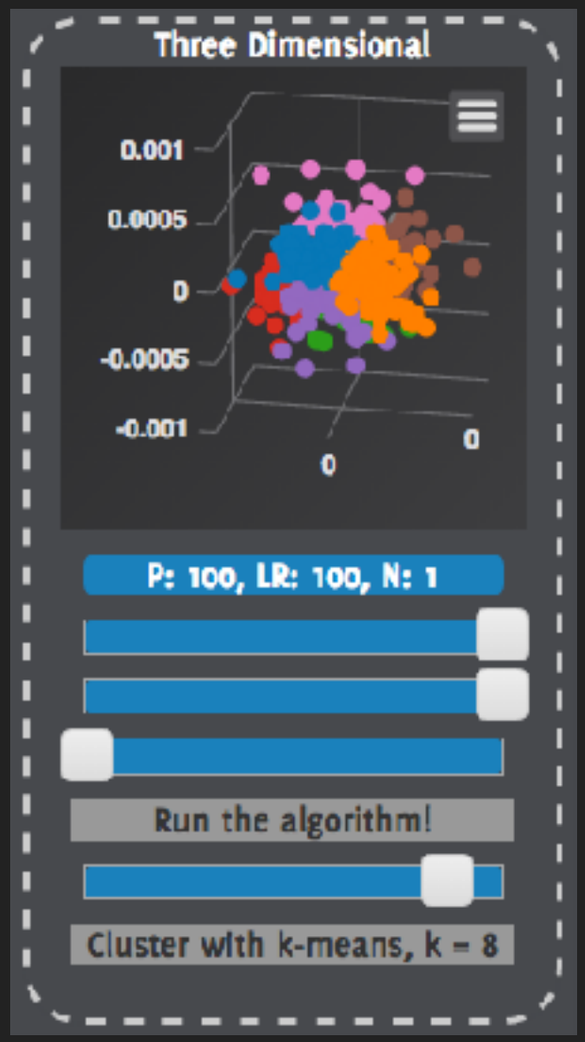

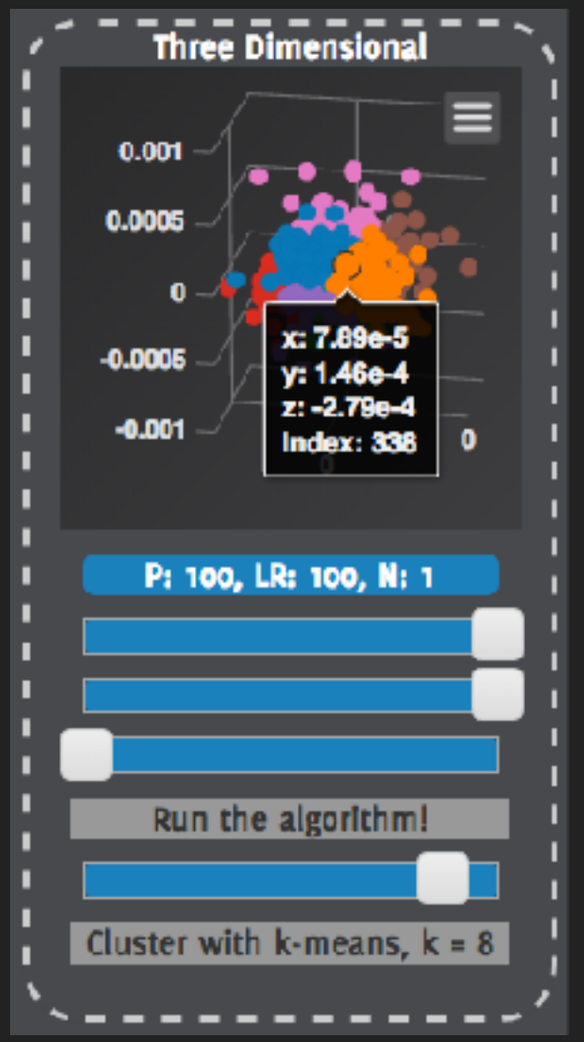

- Users can cluster using k-means in 3 dimensions, then pan and rotate around the data points
- Tooltips appear with information about the point location and its original index in the dataset
- This allows users to identify points that are similar in high dimensions, as candidates for further exploration during feature selection

VISUALIZATION WIDGET: HYPERLINKS TO MORE INFORMATION

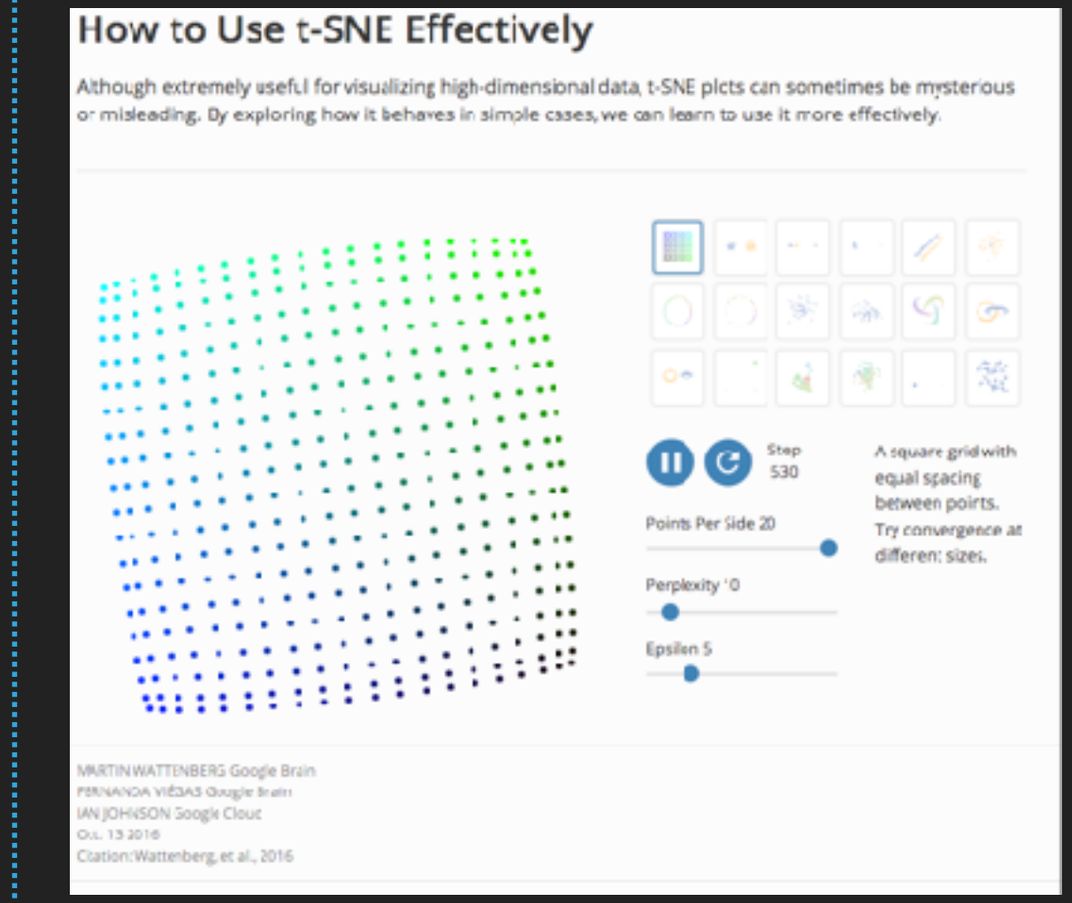

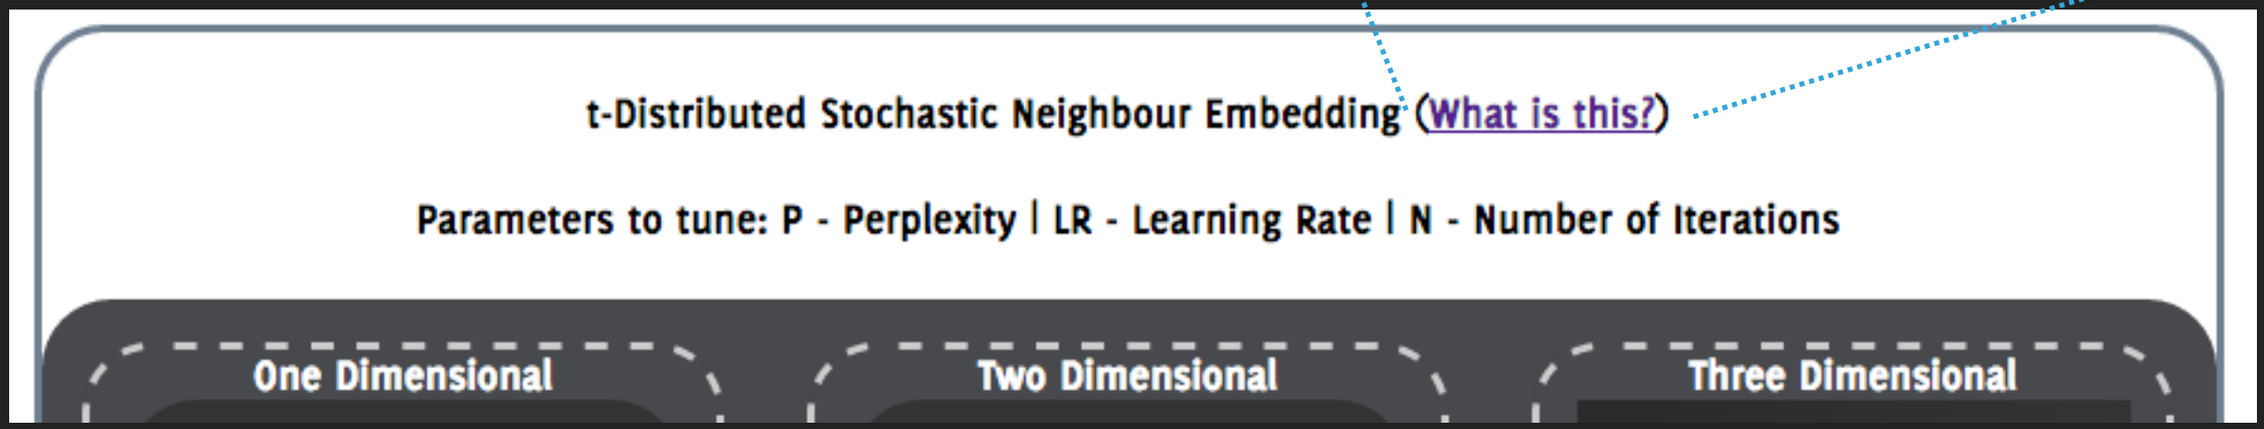

OTHER WIDGETS CURRENTLY IMPLEMENTED: MDS

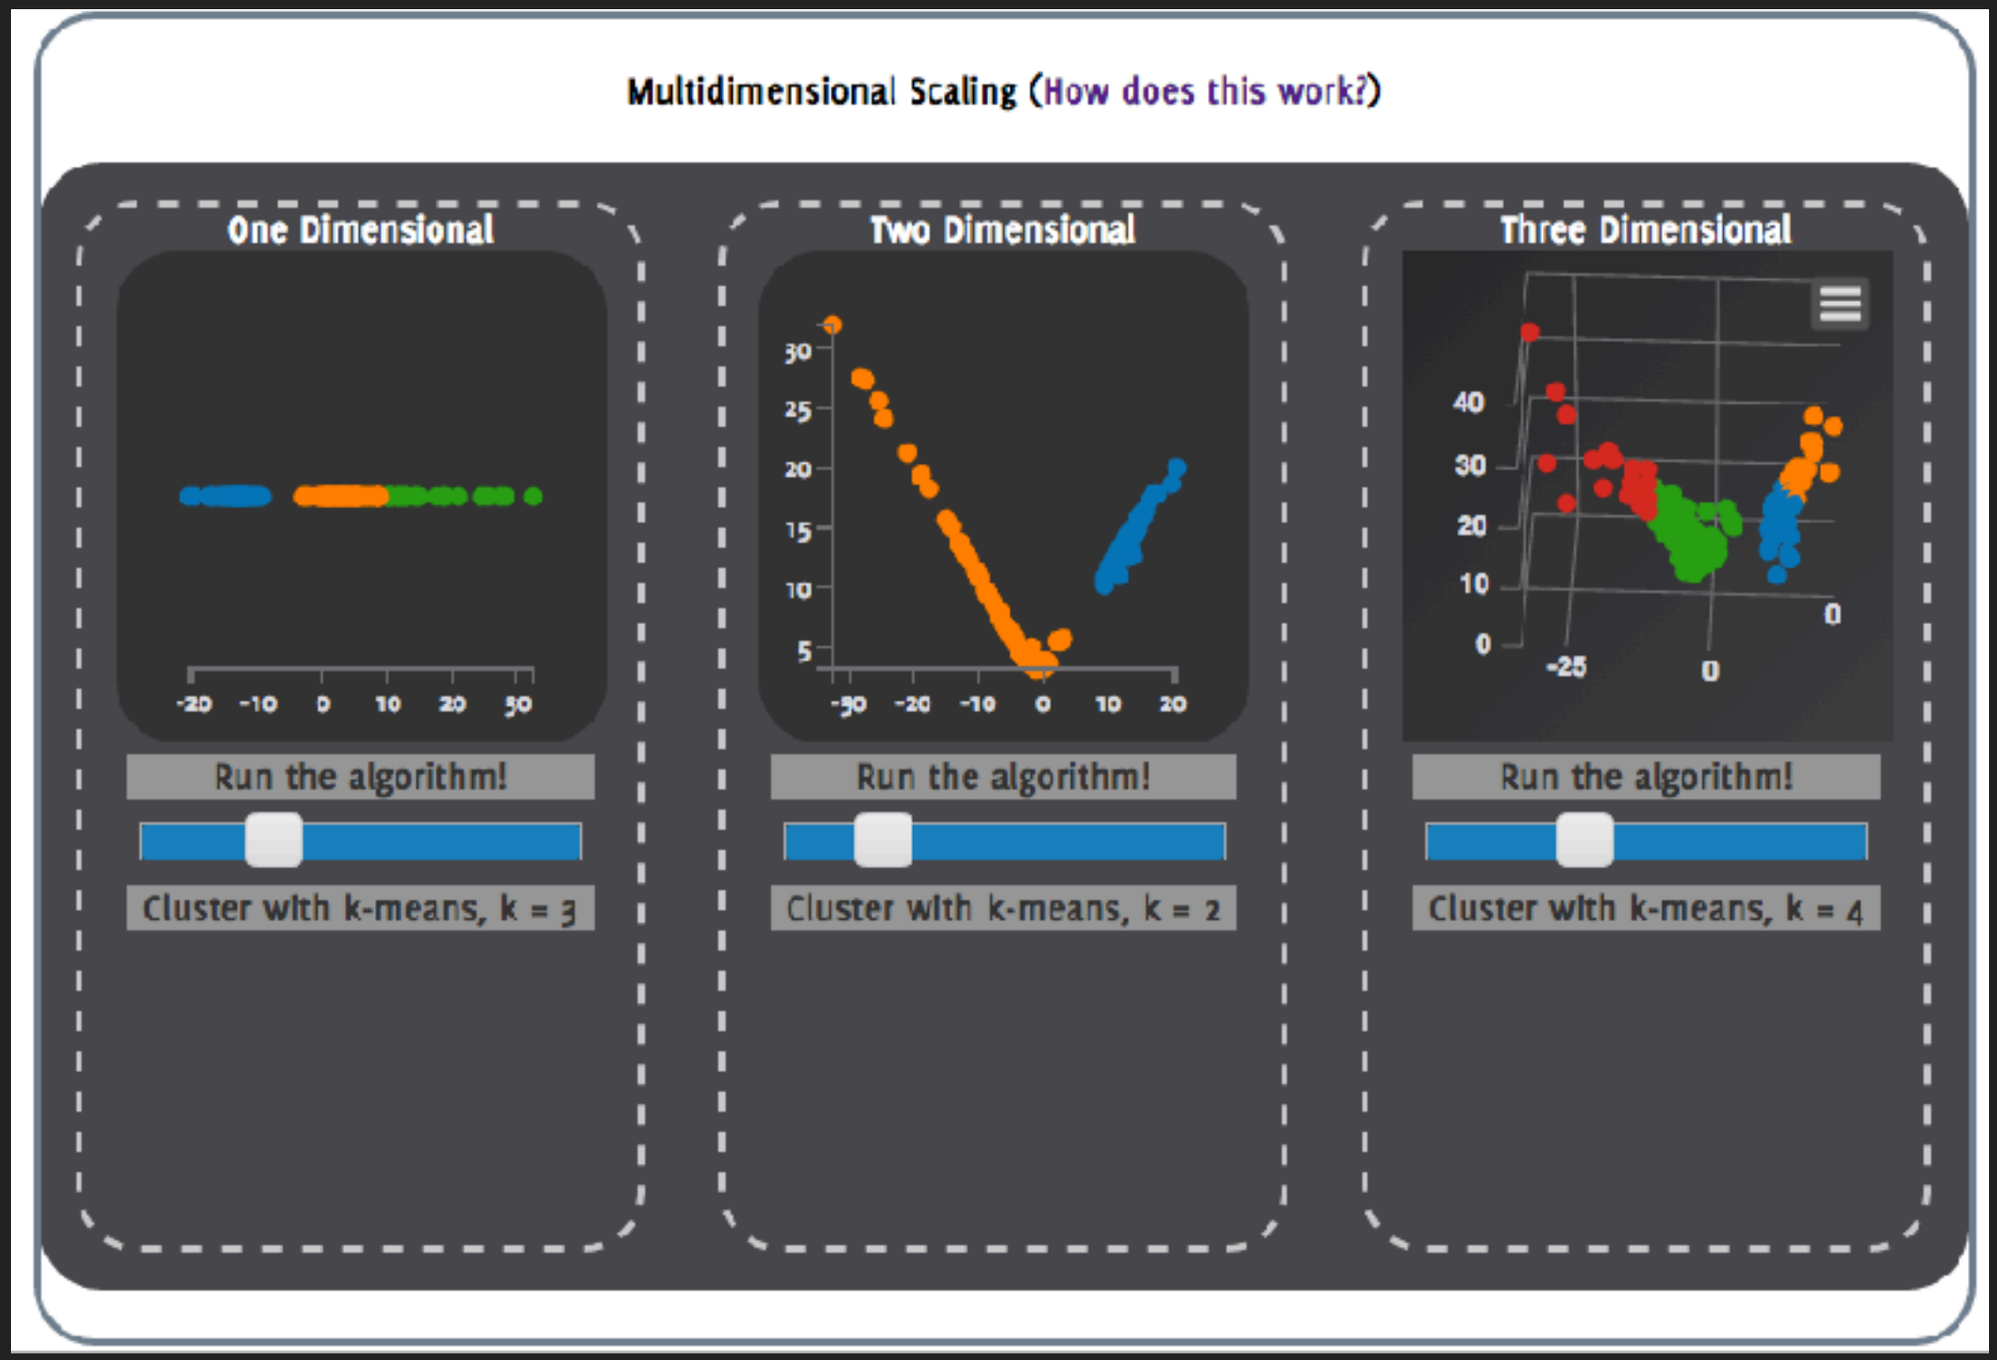

OTHER WIDGETS CURRENTLY IMPLEMENTED: PCA

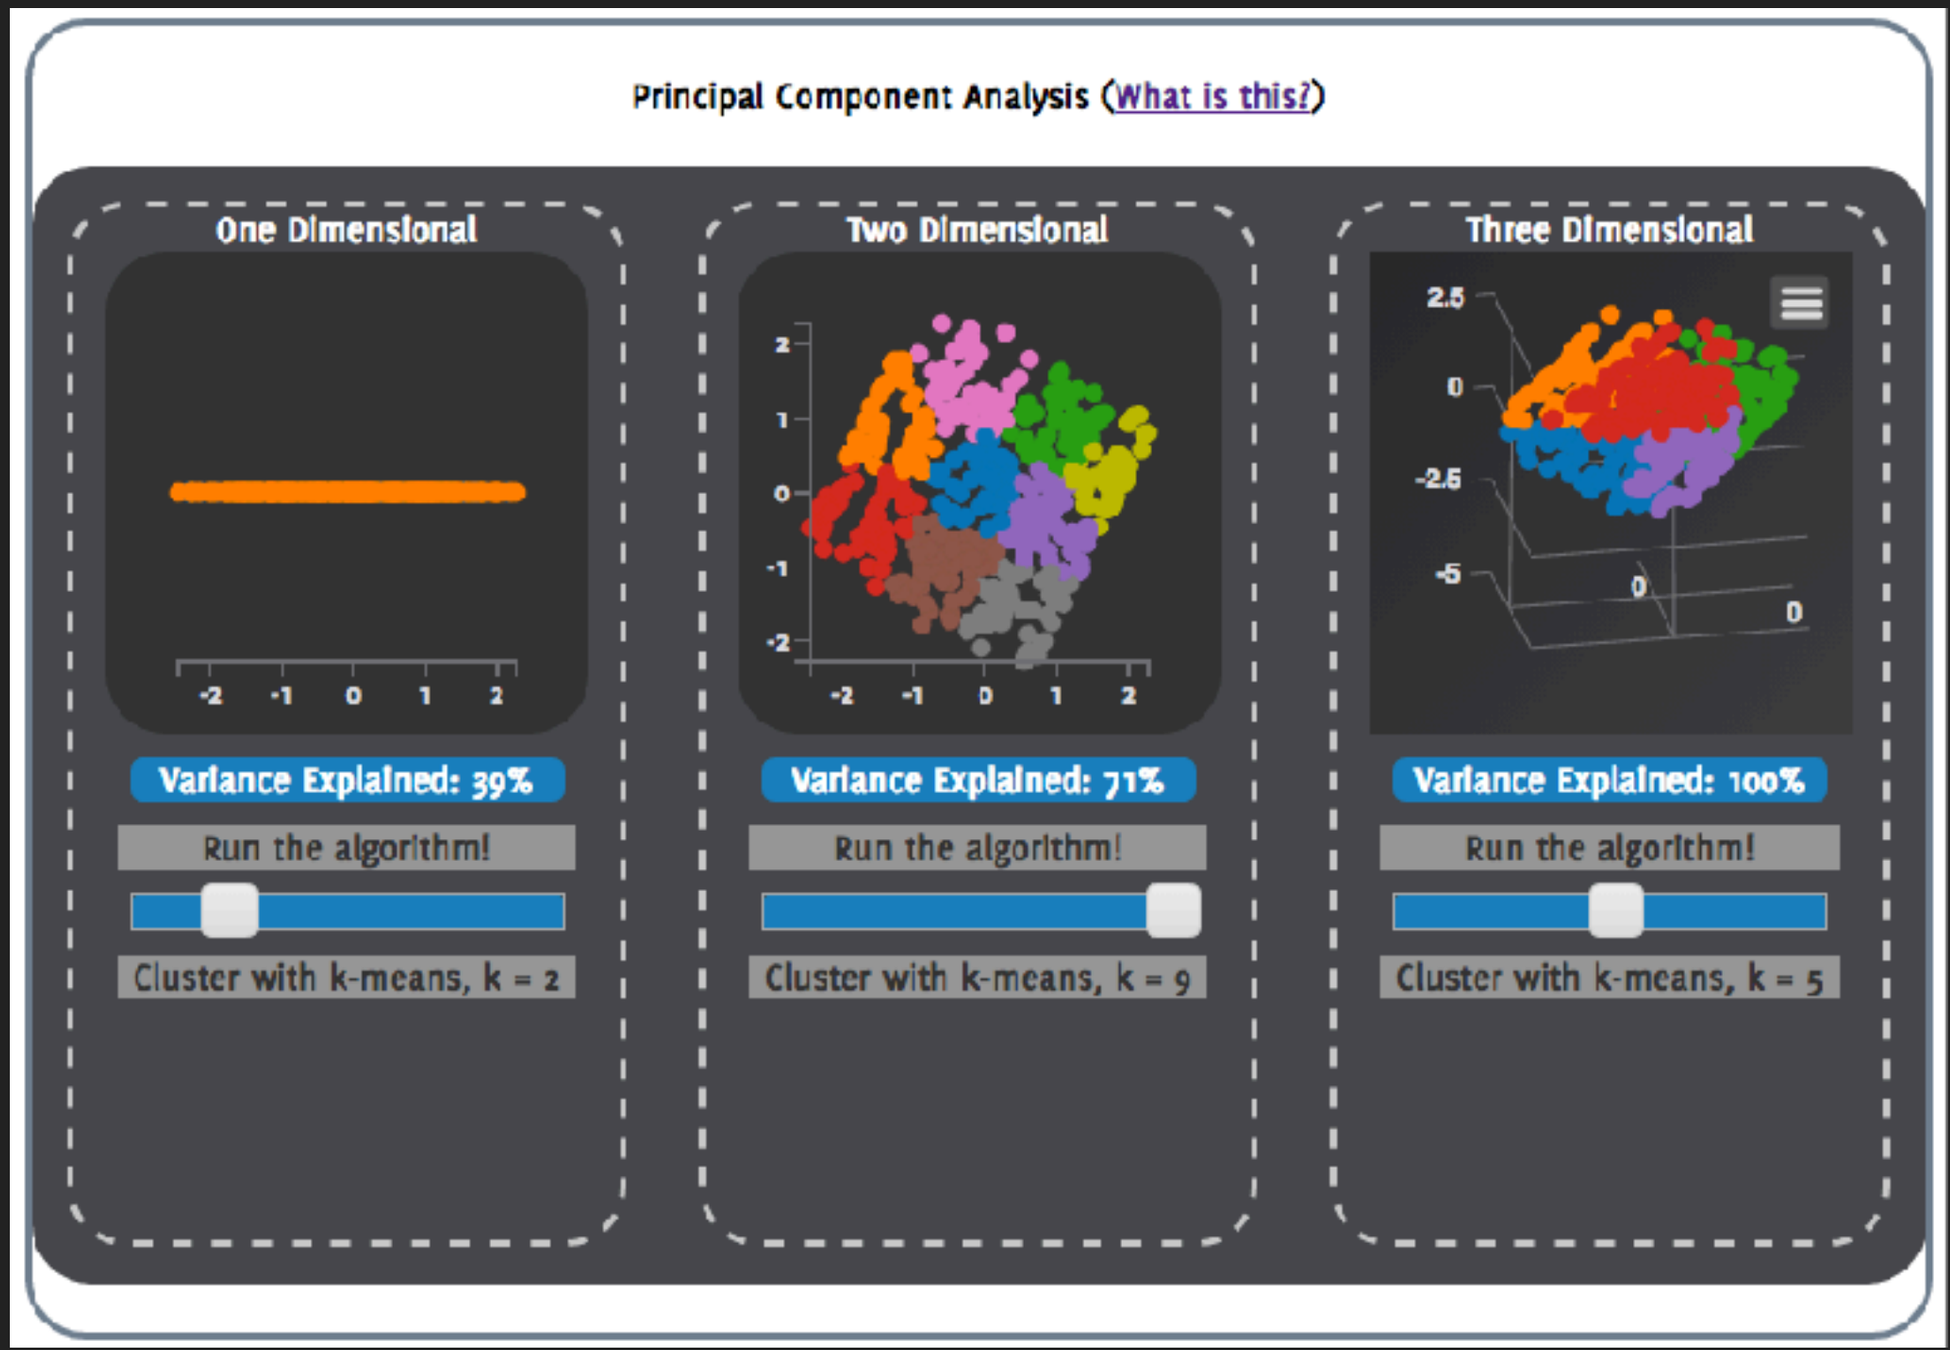

LIVE DEMO: FINDING LATENT STRUCTURE IN DATASET

LIVE DEMO: DISCOVERING CLUSTERS IN FLOWERS DATASET

FORMATIVE STUDY: DESCRIPTION

- One pilot participant, graduate student in Machine Learning group at University of Toronto
- Unstructured interview format
- Given no description about interface and asked to determine its function for studying data stored in a comma-separated file
- Participant discovered data-loading and dimension reduction features by playing with interface

FORMATIVE STUDY: ACTIONABLE FEEDBACK

- Participant requested dynamic data table to edit individual rows and get immediate feedback in visualizations
- Confused by technical terminology of parameters in t-SNE dimensionality reduction technique
- After discovering 3D rotation in visualization widget, participant repeatedly tried to pan 2D visualization by clicking and dragging
- Takeaways: useful to allow “drilling down” by zooming into particular groups data points and panning: create a focus area. Dynamic visualization for adding/removing of data is priority for next iteration

CONCLUSION

- Demonstrated a high-dimensional data visualization dashboard that leverages advanced machine learning to reduce data to 1, 2, or 3 dimensions of variation
- Shown how this can augment data exploration by enhancing latent structure and natural grouping discovery
- Identified areas for future improvement through formative study with pilot participant